

Dive into Conditional Generative Adversarial Networks and Deep Learning

Shaoxuan Chen, Advisor: Yves Atchadé

Abstract

The aim of this report is to summarize the research work that I did in 2020 summer under the guidance of Prof. Yves Atchadé, which basically including two topics: 1. The theory and math derivation of Conditional Generative Adversarial Networks(CGANs) and how it can be applied in regression, i.e., Adversarial Regression; 2. Basic Deep Learning topics that I tackled with, including Linear Neural Networks, Multilayer Perceptrons and basic Deep Learning computation in TensorFlow, Keras. The results of different experiments are also included, like comparing different methods in parameter approximation of linear/logistic regression, binary classification by Linear Neural Network and Multilayer Perceptrons, adjusting and checking the influence of different loss functions and hyperparameters, etc.

1 Conditional GANs

1.1 GANs & Regression

As for regression, if we have the data $\{x_i, y_i\}_{i=1}^n$ following the unknown true joint probability density function(pdf) $P_{data}(x, y)$, the aim of regression is to estimate the unknown true conditional distribution $P(y|x)$ and eventually do prediction if we have the new data came in, i.e., $P(y^*|x^*; x, y)$. However, there were some limitations of regression: 1.The parameters of the model were unknown so there might be too many solutions; 2.The true model is unknown, the underlying stochasticity of the model because of the random noise term ε ; 3. The true model can be any functions even without closed forms. So the number of parameters can be very large, which will prevent the $X^T X$ from invertible.[1]

As for GANs, which contains the Generator and Discriminator. By feeding random noise to the Generator, it can generate the fake data follows an underlying pdf to fool the Discriminator. The Discriminator is fed by data both from Generator, i.e., the fake data, and data from true distribution. So it can be used to estimate the density by measuring the divergence of Generator distribution and true distribution for optimizing the objective function of Generator. Therefore, by combining GANs and Regression, we are also able to estimate the conditional distribution $P(y|x)$, which is also called Adversarial Regression.

1.2 Neural Network

Since GANs are neural network based, so neural network is also briefly mentioned this section. Below, as shown in Figure 1, are the two kinds of version of networks: shallow feedforward neural network and deep neural network.

We can think of neural network as non-linear multi-regression model. And it has three parts: the input layer, hidden layers and output layer. Each layer contains a number of units. For each units there basically happens two operations:

1. Each layer compute a linear combination of the outputs of the previous layer. Note: W represents the weight, b represents the bias, l represents the order of layers

$$z_j^{[l]} = [W_j^{[l]}]^T a^{[l-1]} + b_j^{[l-1]}$$

2. Linear combination $z_j^{[l]}$ transformed by non-linear activation function $g^{[l]}$, which is the output

$$a_j^{[l]} = g(z_j^{[l]}) = g([W_j^{[l]}]^T a^{[l-1]} + b_j^{[l-1]})$$




Figure 1: Simple Version of Neural Network. (reference link: missinglink.ai)

Therefore, combined what have been mentioned above, the whole neural network can be expressed as:

$$A^{[l]} = g(z^{[l]}) = g[(W^{[l]})^T A^{[l-1]} + b^{[l-1]}]$$

The non-linear activation functions allow the model to create complex mappings between the network's inputs and outputs, which are essential for learning and modeling complex data, such as images, video, audio, and data sets which are non-linear or have high dimensionality. [2] The common used non-linear activation functions are Sigmoid, ReLU, Leaky ReLU, Softmax activation functions.

1.3 Theory and math derivation of GANs

The GAN contains a generative model G that captures the data distribution, and a discriminative model D that estimates the probability that a sample came from the training data rather than G . The training procedure for G is to maximize the probability of D making a mistake. This framework corresponds to a minimax two-player game. In the space of arbitrary functions G and D , a unique solution exists, with G recovering the training data distribution and D equal to $\frac{1}{2}$ everywhere. In the case where G and D are defined by multilayer perceptrons, the entire system can be trained with backpropagation. There is no need for any Markov chains or unrolled approximate inference networks during either training or generation of samples.[3]

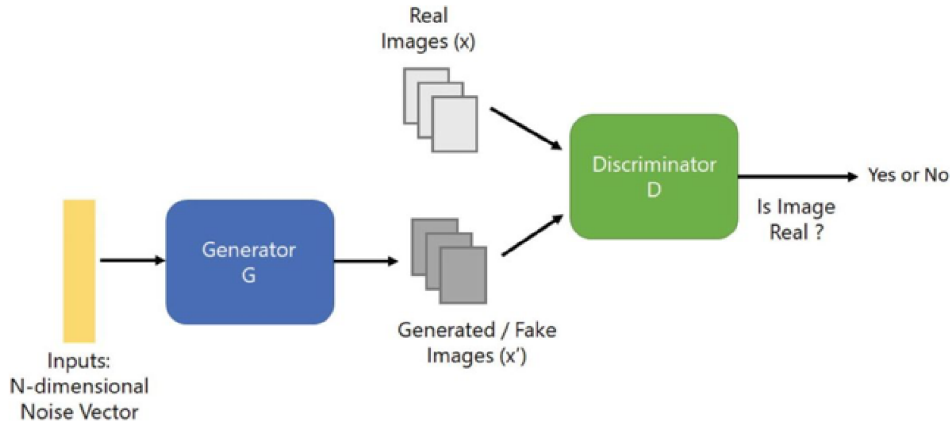


Figure 2: GANs basic framework (reference link: missinglink.ai)

Since the task of Discriminator D is about classification, so the form of the loss function should satisfy Cross Entropy loss function form. To be more specific:

$$V(D, G) = E_{x \sim P_{data(x)}}[\log D(x)] + E_{Z \sim P_{Z(z)}}[\log(1 - D(G(z)))]$$

The adversarial game comes from the Discriminator wants the $D(G(z))$ to decrease. But the Generator wants $D(G(z))$ to go up. Thus, the optimization equation satisfies the form:

$$\min_G \max_D V(D, G) = E_{x \sim P_{data(x)}} [\log D(x)] + E_{Z \sim P_{Z(z)}} [\log(1 - D(G(z)))]$$

Below we want to show that, as for the optimization function, the minimum of cost function is achieved if and only if the probability distribution of Generator match the real data, i.e., the fake sample \approx real sample.

1 > For Discriminator D:

Claim: ① For a fixed Generator, the maximum of cost function w.r.t discriminator is a constant $V(D^*, G) = -\log(4)$

Proof:

$$\begin{aligned} V(D, G) &= E_{x \sim P_{data(x)}} [\log D(x)] + E_{Z \sim P_{Z(z)}} [\log(1 - D(G(z)))] \\ &= \int_x P_{data(x)} \log(D(x)) dx + \int_z P_{Z(z)} \log(1 - D(G(z))) dz \\ &= \int_x P_{data(x)} \log(D(x)) dx + \int_x P_{g(x)} \log(1 - D(x)) dx \\ &= \int_x P_{data(x)} \log(D(x)) + P_{g(x)} \log(1 - D(x)) dx \end{aligned} \tag{1}$$

The inner part satisfies the form of:

$$a \log(y) + b \log(1 - y)$$

If we want to find y to maximize it:

$$\begin{aligned} \frac{\partial a \log(y) + b \log(1 - y)}{\partial y} &= 0 \\ \frac{a}{y} + \frac{b}{1 - y} (-1) &= 0 \end{aligned}$$

Then we have:

$$y = \frac{a}{a + b}$$

Then we plug in the original variable for a and b, we then recover the final form for the optimal discriminator, which can maximize the loss function:

$$D(x)^* = \frac{P_{data(x)}}{P_{data(x)} + P_{g(x)}}$$

For $P_{data(x)} = P_{g(x)}$, the pdf of real data and generated data are identical. So the optimal discriminator returns value:

$$D(x)^* = \frac{P_{data(x)}}{2P_{data(x)}} = \frac{1}{2}$$

Plug in $D(x)^* = \frac{1}{2}$ into our cost function, we then have,

$$\begin{aligned} V(D, G) &= E_{x \sim P_{data(x)}} [\log D(x)] + E_{Z \sim P_{Z(z)}} [\log(1 - D(G(z)))] \\ &= E_{x \sim P_{data(x)}} [\log(\frac{1}{2})] + E_{Z \sim P_{Z(z)}} [\log(\frac{1}{2})] \\ &= -\log(2) - \log(2) \\ &= -\log(4) \end{aligned} \tag{2}$$

Now we have two conclusions: The first one is that the optimal discriminator $D(x)^* = \frac{P_{data(x)}}{P_{data(x)} + P_{g(x)}}$; The second one is that when $P_{data(x)} = P_{g(x)}$, then we have $D(x)^* = \frac{1}{2}$, $V(D^*, G) = -\log(4)$. That is

to say, when distribution is equal, the upper bound of cost function is $-\log(4)$.

2 > For Generator G:

For Generator, we want to minimize the cost function $V(D, G)$. Here we need to introduce the basic concept of Kullback-Leibler(KL) divergence and Jensen-Shannon(JS) divergence, which both of them represent the distance measure between two probability distributions.

KL Divergence:

$$D_{KL}(P||Q) = E_{x \sim P}[\log \frac{P(x)}{Q(x)}]$$

JS Divergence:

$$JSD(P||Q) = \frac{1}{2}D_{KL}(P||\frac{P+Q}{2}) + \frac{1}{2}D_{KL}(Q||\frac{P+Q}{2})$$

Note that P, Q represents two probability distributions. The JS Divergence is similar to KL Divergence, except that it is symmetric, which means that the JS Divergence from P to Q is the same as the JS Divergence from Q to P. And this property does not hold in KL Divergence.

Notice that the form of KL Divergence is exactly the same as if we plug in the expression of optimal Discriminator, i.e.,

$$\begin{aligned} D_{KL}(P||Q) &= E_{x \sim P}[\log \frac{P(x)}{Q(x)}] \\ &= E_{x \sim P}[\log \frac{\frac{1}{2}P(x)}{\frac{1}{2}Q(x)}] \\ &= E_{x \sim P}[\log \frac{P(x)}{Q(x)/2}] - \log(2) \end{aligned} \tag{3}$$

Then we rewrite the cost function, we have:

$$V(D^*, G) = -\log(4) + KL(P_{data}||\frac{P_{data} + P_g}{2}) + KL(P_g||\frac{P_{data} + P_g}{2})$$

Combined with the form of JS Divergence, we have:

$$V(D^*, G) = -\log(4) + 2JSD(P_{data}||P_g)$$

Since the minimum of any JS Divergence is 0, and it occurs if and only if $P_{data(x)} = P_{g(x)}$, then for the objective function with optimal Discriminator plugged in, we have the minimum of the function:

$$\min_G V(D^*, G) = -\log(4)$$

So we know that the objective function $V(D^*, G)$ has one minimum $-\log(4)$, and the minimum is achieved when generator perfectly match the real data distribution.

1.4 Conditional GANs Algorithm

The framework of Conditional GANs is almost the same as the basic GANs. The only difference is that the input is the data together with labels instead of just data. More details of comparison between the objective function are shown below, where y represent the label:

GAN:

$$\min_G \max_D V(D, G) = E_{x \sim P_{data(x)}}[\log D(x)] + E_{Z \sim P_{Z(z)}}[\log(1 - D(G(z)))]$$

CGAN:

$$\min_G \max_D V(D, G) = E_{x \sim P_{data(x)}}[\log D(x|y)] + E_{Z \sim P_{Z(z)}}[\log(1 - D(G(z|y)))]$$

Below in Figure 3, which is cited in paper Conditional Generative Adversarial Nets[4], the framework of Conditional GAN is shown clearly.

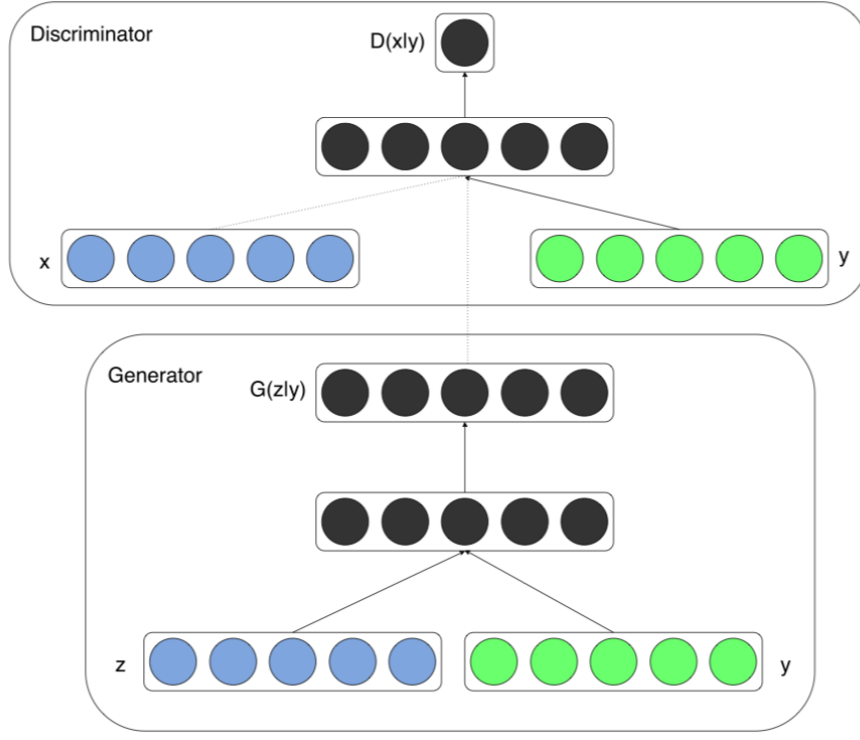


Figure 3: Conditional GANs basic framework[4]

As for Adversarial Regression, the algorithm of Conditional GANs for Regression is shown below:

Algorithm 1: Conditional GANs Algorithm for regression

Required: M the minibatch size. A prior function generating z . The hyperparameters d_{steps} and g_{steps} . An algorithm for gradient ascent. The number of iterations.

```

1 for number of iterations do
2   for  $d_{steps}$  do
3     Sample minibatch of labels  $x_1, \dots, x_m$  from data  $p_{data}(x) \sim \text{Uniform}(0,1)$ ;
4     Produce sample data  $y_1, \dots, y_m$  by corresponding  $x_1, \dots, x_m$  via the regression model,
       label with 1;
5     Sample minibatch of examples  $z_1, \dots, z_m$  from noise prior  $p_g(z)$ , transform with
       Generator to have  $\{G(z_i)\}_{i=1}^m$ , label with 0;
6     Past the sample data and labels  $\{x_i\}_{i=1}^m$  to Discriminator to get predictions
        $\{D(y_i | x_i)\}_{i=1}^m$  &  $\{D(G(z_i | x_i))\}_{i=1}^m$ ;
7     Update  $\theta_D$  by ascending:

```

$$\nabla_{\theta_D} \frac{1}{m} \sum_{i=1}^m \log(D(y_i | x_i)) + \log(1 - D(G(z_i | x_i)))$$

```

8   end
9   for  $g_{steps}$  do
10    Sample minibatch of examples  $z_1, \dots, z_m$  from noise prior  $p_g(z)$ , transform with
       Generator to have  $\{G(z_i | x_i)\}_{i=1}^m$ ;
11    Update  $\theta_G$  by ascending the non-saturating function:

```

$$\nabla_{\theta_G} \frac{1}{m} \sum_{i=1}^m \log(D(z_i | x_i))$$

```

12  end
13 end

```

2 Computation Experiments

This section relates to the directed reading I did with Prof. Yves Atchadé of the e-book Dive into Deep Learning and the application of the knowledge I learned. It mainly covers the experiments I did related to parameter approximation by Mini-batch Stochastic Gradient Descent(SGD)/Stochastic Gradient Descent/analytic approach, and the experiments related to binary classification by Logistic Regression(Linear Neural Network) and Multilayer Perceptrons(MLP). Results of different hyperparameters were compared like learning rate, batch size, number of epochs. The performance of different methods were also compared in dealing with binary classification, like with/without L2 penalty terms in loss function, coding manually from scratch, concise implementation in Keras, etc.

2.1 Optimizer

Algorithm for SGD

Algorithm 2: SGD

Required: The loss function l . The number epochs. Size of data \mathcal{N} , learning rate η

- 1 Initialize parameters (\mathbf{w}, b)
- 2 **for** *number of epochs* **do**
- 3 Compute gradient \mathbf{g} on the whole data set $\leftarrow \partial_{(\mathbf{w}, b)} \sum_{i \in \mathcal{N}} l(\mathbf{x}^{(i)}, y^{(i)}, \mathbf{w}, b)$
- 4 Update parameters $(\mathbf{w}, b) \leftarrow (\mathbf{w}, b) - \eta \mathbf{g}$
- 5 **end**

Algorithm for Mini-batch SGD

Algorithm 3: Mini-batch SGD

Required: \mathcal{B} the minibatch size. The loss function l . The number epochs. learning rate η

- 1 Initialize parameters (\mathbf{w}, b)
- 2 Shuffle the data based on the batch size
- 3 **for** *number of epochs* **do**
- 4 Compute gradient \mathbf{g} based on the shuffled data $\leftarrow \partial_{(\mathbf{w}, b)} \frac{1}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} l(\mathbf{x}^{(i)}, y^{(i)}, \mathbf{w}, b)$
- 5 Update parameters $(\mathbf{w}, b) \leftarrow (\mathbf{w}, b) - \eta \mathbf{g}$
- 6 **end**

2.2 Parameter Approximation of Linear Regression by Linear Neural Network

Model for generating the true data:

$$y = XW + b + \varepsilon, \text{ where } X_i^\top \sim \mathcal{N}\left(\begin{pmatrix} 0 \\ 0 \end{pmatrix}, \text{diag}(1, 1)\right), W = [2, 3.5]^\top, b = 4, \varepsilon_i \sim N(0, 0.01)$$

Loss Function: Squared Loss Function

$$L = \frac{1}{m} \sum_{i=1}^m (\hat{y}_i - y_i)^2$$

The results of different approaches/hyperparameters in parameter approximation of Linear Regression are shown in Table 1.

From below, as for Mini-batch SGD, we can see that for small fixed batch size, if the learning rate is too big, i.e., equals to 1. Then after the same number of epochs, the one with smaller learning rate have better performance than the one with larger learning rate. That may because that larger learning rate takes larger steps in finding the optimal point. If the learning rate is too big, instead of approaching the optimal point step by step, the gradient descent approach will bouncing around the optimal point and never reach the optimal point. Also, as for batch size in mini-batch SGD, we can see that it is better to use larger learning rate if the batch size is large, and smaller learning rate if the batch size is small. As for SGD, since it use the whole data set to do optimization, the result shows that larger learning have better performance. Also, comparing Mini-batch SGD with SGD, we can see

Table 1: Results of different approaches in Linear Regression parameter approximation(Note: $\sigma^2 = 1$)

Optimizer	lr	epochs	batch size	error in estimating w & b
Mini-batch SGD	0.03	5	10	[-0.00027895 -0.00031495], [0.00035882]
Mini-batch SGD	1	5	10	[0.00152123 0.00266123], [0.00423193]
Mini-batch SGD	0.03	5	50	[0.06462872 0.1595416], [0.16159701]
Mini-batch SGD	1	5	50	[-0.0001142 0.00045729], [-0.00081253]
SGD	0.03	5	\	[1.6979696 3.0017211], [3.4078]
SGD	0.5	5	\	[0.03197384 0.10298157], [0.09555578]
SGD	1	5	\	[-4.9352646e-05 -7.7486038e-05], [2.360344e-05]
Analytic approach	\	\	\	[-6.918907e-04 -4.553795e-05],[4.220009e-05]

that after adjustment, they can have almost the same performance in parameter approximation. The analytic approach also has very good performance.

2.3 Parameter Approximation of Logistic Regression by Linear Neural Network

Model for generating the true data:

$$\log \frac{p}{1-p} = XW + b, \text{ where } X_i^\top \sim \mathcal{N}\left(\begin{pmatrix} 0 \\ 0 \end{pmatrix}, \text{diag}(1, 1)\right), W = [2, 3.5]^\top, b = 4$$

Activation Function: Sigmoid Function

$$h(x) = \sigma(wx + b), \text{ where } \sigma(z) = \frac{1}{1 + e^{-z}}$$

Loss Function: Cross Entropy Loss Function

$$L = -\frac{1}{m} \sum_{i=1}^m y_i \log a_i + (1 - y_i) \log(1 - a_i)$$

$$a_i = \text{sigmoid}(z_i), \text{ where } z_i = W^T x_i + b$$

Loss Function with L_2 penalty term:

$$L = -\frac{1}{m} \sum_{i=1}^m y_i \log a_i + (1 - y_i) \log(1 - a_i) + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

$$a_i = \text{sigmoid}(z_i), \text{ where } z_i = W^T x_i + b$$

Refined Loss Function built in TensorFlow: `tf.nn.softmax_cross_entropy_with_logits`

$$L = \max(a_i, 0) - a_i * y + \log(1 + \exp(-\text{abs}(a_i)))$$

$$a_i = \text{sigmoid}(z_i), \text{ where } z_i = W^T x_i + b$$

2.3.1 Results of using Cross Entropy Loss Function

The results of different approaches/hyperparameters in parameter approximation of Logistic Regression are shown in Table 2.

From the results shown below, we can see that Mini-batch SGD approach works well when the learning rate is relatively small. As the learning rate increase, even though the rate of converging to true values increase very fast, i.e., the true value of parameters are estimated relative well even at the first epoch, the value of the estimated parameters also fluctuate a lot. Since the objective function is a convex function, so the value of learning rate does not have too much effect. However, tuning the learning rate also make sure that we can have better results. Also, we notice that if we increase the

Table 2: Results of different approaches in Logistic Regression parameter approximation(Note: $\sigma^2 = 1$)

Optimizer	lr	epochs	batch size	error in estimating w & b
Mini-batch SGD	0.1	10	10	[0.04171479 -0.02211356], [0.08550978]
Mini-batch SGD	0.5	10	10	[0.08610749 -0.09116745], [-0.20888281]
Mini-batch SGD	1	10	10	[-0.03445697 0.09493375], [-0.39592457]
Mini-batch SGD	0.1	10	50	[0.37641394 -0.58421636], [0.65578675]
Mini-batch SGD	1	10	50	[0.01389563 -0.0771153], [-0.0222297]
SGD	0.01	100	\	[1.9053147 -3.225234], [3.9055605]
SGD	0.1	100	\	[1.4485135 -2.4557726], [2.7461467]
SGD	0.5	100	\	[0.82667005 -1.3771112], [1.5490735]
SGD	1	100	\	[0.5369544 -0.876415], [0.99448943]
SGD	1	500	\	[0.09156346 -0.10711193], [0.12568903]
SGD	1	1000	\	[0.05528092 -0.04446769], [0.05429029]

batch size, the one with larger learning rate have better performance. As for SGD approach, larger learning rate have better performance. And even using learning rate = 1, the estimated parameters converge very slow at the later steps.

However, as shown in Table 3, if we change the σ^2 to 100 instead of 1 when generating X. Neither the Mini-batch SGD nor the SGD approach work anymore.

Table 3: Results of different approaches in Logistic Regression parameter approximation(Note: $\sigma^2 = 100$)

Optimizer	lr	epochs	batch size	error in estimating w & b
Mini-batch SGD	0.1	10	10	[nan nan], [nan]
Mini-batch SGD	1	10	10	[nan nan], [nan]
SGD	0.1	100	\	[nan nan], [nan]
SGD	0.5	100	\	[nan nan], [nan]
SGD	1	100	\	[nan nan], [nan]

The reason of that may because the large variance when generating the data may cause the $1 - a_i$ term, i.e., parts including sigmoid function, in the loss function equals to zero, so the $\log(1 - a_i)$ term will blow up to infinite, and the gradient of it will become extremely small, which is the problem of vanishing gradient. Below in Figure 4 also clearly show the rationale of that. So we can conclude that the non-refined loss function, i.e., Cross Entropy Loss Function, is not robust enough to deal with all kinds of data.

2.3.2 Results of using L_2 penalty terms in loss function.

The result of including L_2 penalty term in loss function is a standard techniques for regularizing models. However, the result of adding penalty terms in these experiments did not have better performance compared to no penalty term in loss function. So the penalty term might be better used in higher dimension data.

2.3.3 Results of using TensorFlow Built in Loss Function

The results of different approaches/hyperparameters in parameter approximation of Logistic Regression using the TensorFlow built in loss function, i.e., *tf.nn.softmax_cross_entropy_with_logits*, are shown in Table 4.

The results shown that if we use the refined loss function, i.e., the built in cross entropy loss function in TensorFlow, after tuning the hyperparameters, the true value of parameters can be approached at certain epochs, e.g., the Mini-batch SGD approach at epoch 2, SGD approach at epoch 1000. However,




Figure 4: Sigmoid Function and the Gradient (reference link: Vanishing Gradient with Sigmoid activation in dnns)

Table 4: Results of different approaches in Logistic Regression parameter approximation(Note: $\sigma^2 = 1$)

Optimizer	lr	epochs	batch size	error in estimating w & b
Mini-batch SGD	0.5	2	10	[-0.1762023 0.4935906], [-0.26229382]
Mini-batch SGD	0.5	5	10	[-1.0851371 1.7294331], [-1.6065059]
Mini-batch SGD	0.5	10	10	[-1.7718003 2.7897825], [-2.9039044]
SGD	0.5	100	\	[1.2511575 -2.1367278], [2.7508674]
SGD	0.5	1000	\	[0.02554774 -0.05595708], [0.46379805]
SGD	0.5	1000	\	[-0.48387456 0.7847338], [-0.47854042]

in stead of converging to the true parameter, the estimated value of parameters continue increase or decrease as the number of epochs increase and they never converge. Also, the same experiments were also did if $\sigma^2 = 100$. The results shown that the built in loss function in TensorFlow can deal with data with large variance, i.e., no more vanishing gradient problems shown in the results. But the estimated results are also bad. One of the reason might be the true loss function is fundamentally changed to deal with the vanishing gradient problem. Thus, the optimal point of the objective function is also changed, which means the parameters we used in generating the true data are not the optimal parameters for the result anymore. So it is a more robust loss function to deal with data with more randomness at the expense of the accuracy.

2.4 Binary Classification by Linear Neural Network & MLP

Last section focus primarily on parameter approximation. This section focus on training and testing the model. The data was spilt into training & testing data set(note: training v.s testing = 4:1), and the model was trained on the training data set by Linear Neural Network/MLP approach and tested on the testing data. The results after certain number of epochs are shown in Table 5. And the plot of the result is also shown in Figure 5.

Linear Neural Network: define above in 2.3, using the Sigmoid Function as the activation function

Size of the data: 1000

Non-Linear Activation Function in MLP: Rectified Linear Unit(ReLU)

$$ReLU(x) = \max(0, x)$$

MLP Model:

$$H = \sigma(XW^{(1)} + b^{(1)})$$

$$O = HW^{(2)} + b^{(2)}$$

Note: σ : the non-linear activation function; $X \in R^{n \times d}$, $H \in R^{n \times h}$, hidden-layer weights $W^{(1)} \in R^{n \times d}$, out-put layer weights $W^{(2)} \in R^{h \times q}$, biases $b^{(1)} \in R^{1 \times h}$, biases $b^{(2)} \in R^{1 \times q}$; d:number of inputs; h:number of hidden units; q:number of outputs; H:Hidden Layer; O: Output Layer

Loss Function: Refined Loss Function built in TensorFlow(defined above in 2.3)

Number of Hidden Layer, Output Layer: 1, 1

Number of Hidden Units: 256

2.4.1 Experiments did on data generated by Logistic Regression

Model for generating the true data:

$$\log \frac{p}{1-p} = XW + b, \text{ where } X_i^\top \sim \mathcal{N}\left(\begin{pmatrix} 0 \\ 0 \end{pmatrix}, \text{diag}(1,1)\right), W = [2, 3.5]^\top, b = 4$$

The experiment was done by mainly two computing approaches: coding manually from scratch v.s. concise implementation in Keras. Both of the results are shown in Table 5. And the training, testing processes are also shown in Figure 5 & 6. From the results shown below, we can see that after certain number of epochs, the model was trained very well and the test accuracy reach up to 0.87. Among different approaches, we can see that the performance of them are almost the same. All of them have good test accuracy. Moreover, from the training and testing plot shown in Figure 5 and 6. Note that the train loss was scaled by 35 to fit in the plots. It is more clear in Figure 6 that there is a big jump at epoch 6, which is the reflection of the gradient descent process.

Table 5: Results of different approaches in Binary Classification(Note: $\sigma^2 = 1$, LNN:Linear Neural Network, acc: accuracy, ns: not shown; manual: coding manually from scratch; Keras: concisely implemented by Keras)

Method	Optimizer	lr	epochs	batch size	train loss	train acc	test acc
LNN-manual	Mini-batch SGD	0.01	15	32	16.355714	0.8475	0.8705
MLP-manual	Mini-batch SGD	0.01	30	32	16.649528	0.8475	0.8705
MLP-Keras	Mini-batch SGD	0.01	10	32	16.655714	ns	0.8612




Figure 5: LNN-manual training/testing process




Figure 6: MLP-manual training/testing process

2.4.2 Experiments did on NBA data

The methods shown above were also applied in real life data—the 5-Year Career Longevity for NBA Rookies data, which includes 1340 entries and 19 features like Games Played, Points Per Game, 3 Point Attempts. The response variable is whether the career of the rookies is greater or equal to five years. The data was also centered and normalized beforehand. The results of different approaches are

shown in Table 6. And the training, testing processes are also shown in Figure 7 & 8. Note that the train loss was also scaled by 35 to fit in the plots here.

From the results, we can see that after tuning the hyperparameters, all of the three approaches will eventually reach to almost the same test accuracy. However, overall the test accuracy of three approaches are relative low compare to the experiments did above. So real life data is relatively hard to handle with. From Figure 7 we notice that the test accuracy reach to relative high at the beginning and increase very slow later. The MLP method in Figure 8 shows clearly the increase of test accuracy.

Table 6: Results of different approaches in Binary Classification(Note: $\sigma^2 = 1$, LNN:Linear Neural Network, acc: accuracy, ns: not shown; manual: coding manually from scratch; Keras: concisely implemented by Keras)

Method	Optimizer	lr	epochs	batch size	train loss	train acc	test acc
LNN-manual	Mini-batch SGD	0.001	10	64	17.035067	0.6523	0.6271
MLP-manual	Mini-batch SGD	0.001	10	32	18.476337	0.6544	0.6285
MLP-Keras	Mini-batch SGD	0.001	10	32	16.655714	ns	0.6230




Figure 7: LNN-manual training/testing process




Figure 8: MLP-manual training/testing process

3 Discussion

As for the first part of this report, it is not hard to understand the math derivation of GANs. However, to really do some computing experiments related to GANs, or Adversarial Regression mentioned above, a bunch of computing background knowledge and Deep Learning topics should be mastered with before truly tackle with GANs.

As for the computing experiments did in the later part of the report. It is clear that the hyperparameters have big effects on the training/testing results. So tuning the hyperparameter is an very important task in Deep Learning model building. And if we tune the hyperparameter correctly, our models can reach the same accuracy via any methods. Also, there is still one thing which may need to put more effort which is the higher dimension data. In the experiments shown above, the dimension of data was small and all the approaches work well. But whether those approaches still work well need more systematically experiments to be done. Moreover, for binary classification, I still think the Cross Entropy Loss Function define in TensorFlow is not good enough. Even though it is robust enough to deal with any kind of data, the testing accuracy of that approach is relatively low. The data sets of this project are also relatively easy to deal with. More complicated experiments can be done in the future like dealing with image data.

Overall, this summer project gave me the sense of basic Deep Learning computation in Python and provided me with the chance to get to know the GANs. It also helped me realized the limitation of my

current knowledge and push me to learn more to solid my background to tackle with more interesting problems!

References

- [1] Yoann, B. Adversarial Regression. Generative Adversarial Networks for Non-Linear Regression: Theory and Assessment. ArXiv e-prints, Oct 2019
- [2] Website: missinglink.ai. 7 Types of Neural Network Activation Functions: How to Choose?
- [3] Goodfellow, I. J., Pouget-Abadie, J., Mirza, M., Xu, B., Warde-Farley, D., Ozair, S., Courville, A., and Bengio, Y. Generative Adversarial Networks. ArXiv e-prints, June 2014
- [4] Mirza, Mehdi Osindero, Simon. (2014). Conditional Generative Adversarial Nets.

Code in Python:

```
1  """#linear Regression
2  %matplotlib inline
3  from d2l import TensorFlow as d2l
4  import TensorFlow as tf
5  import random
6  import numpy as np
7  import pandas as pd
8  import TensorFlow_probability as tfp
9  tfd = tfp.distributions
10 from sklearn.model_selection import train_test_split
11
12 num_classes = 1
13 num_features = 2
14 #data generating & processing
15 def logit_data(w, b, num_examples):
16     """Generate  $y = Xw + b$ ."""
17     X = tf.zeros((num_examples, w.shape[0]))
18     X += tf.random.normal(shape=X.shape, stddev=1)
19     y = tf.matmul(X, tf.reshape(w, (-1, 1))) + b
20     #y += tf.random.normal(shape=y.shape, stddev=0.01)
21     y = tf.reshape(y, (-1, 1))
22     return X, y
23
24 true_w = tf.constant([2, -3.4])
25 true_b = 4.2
26 x, y_in = logit_data(true_w, true_b, 10000)
27 p = 1/(1+tf.exp(-y_in))
28 bernoulli_distribution = tfd.Bernoulli(probs=p)
29 X=x
30 sample = bernoulli_distribution.sample(1)
31 y = sample[0] #choose 2 dimensions out of 3 in tensor
32 y = tf.dtypes.cast(y, tf.float32)
33
34 #data processing
35 #split data into training & testing training: testing = 4:1
36 X_train, X_test, y_train, y_test = X[0:8000,:], X[8000:10000,:], y[0:8000,:], y[8000:10000,:]
37 # Feature Matrix # Data labels
38 #X_test.shape, y_train.shape
39
40
41 #shuffle the data
42 def data_iter(batch_size, features, labels):
43     num_examples = len(features)
44     indices = list(range(num_examples))
45     # The examples are read at random, in no particular order
46     random.shuffle(indices)
47     for i in range(0, num_examples, batch_size):
48         j = tf.constant(indices[i: min(i + batch_size, num_examples)])
49         yield tf.gather(features, j), tf.gather(labels, j)
50
51
52 #initialization
53 w = tf.Variable(tf.random.normal(shape=(num_features, num_classes), mean=0, stddev=0.01),
54                 trainable=True)
55 b = tf.Variable(tf.zeros([num_classes]), trainable=True)
56 w,b
57 #activation function
58 def sigmoid(X):
59     return 1 / (1 + np.exp(-X)) #note: use tf.sigmoid later
60 #model
61 def logistic(X, w, b):
62     return tf.sigmoid(tf.matmul(X, w) + b)
63                                     # -1: keep block, use W.shape[0] to calculate left dimension
64 #net(X_train, w, b)
65
66 #loss function
67 def loss(y_hat, y):
68     return tf.reduce_mean(tf.nn.sigmoid_cross_entropy_with_logits(logits=y_hat, labels=y))
69
70 def loss1(y_hat, y):
71     return tf.nn.sigmoid_cross_entropy_with_logits(logits=y_hat, labels=y)
72
73 def loss2(y_hat, y):
74     return tf.losses.binary_crossentropy(
75         y, y_hat, from_logits=True)
```

```

76
77 #loss function---manual
78 #def loss_manual(y_hat, y):
79 #     return -(tf.reduce_sum(y_train*tf.math.log(logistic(X_train, w, b)) +
80 #                             (1-y_train)*(tf.math.log(1 - logistic(X_train, w, b))))) /X_train
81 #     .shape[0]
82 #def loss_manual(y_hat, y, batch_size):
83 #     return -(tf.reduce_sum(y*tf.math.log(y_hat) +
84 #                             (1-y)*(tf.math.log(1 - y_hat))))) /batch_size
85 #def loss_manual_indi(y_hat, y):
86 #     return -(y*tf.math.log(y_hat) + (1-y)*(tf.math.log(1 - y_hat)))
87
88 def loss_check(X, w, b, y):
89     return (tf.matmul(X, w) + b)-(tf.matmul(X, w) + b)*y + tf.math.log(1+tf.math.exp(-(tf.
90     matmul(X, w) + b)))
91
92 def loss_check2(X, w, b, y):
93     return -(tf.matmul(X, w) + b)*y + tf.math.log(1+tf.math.exp(tf.matmul(X, w) + b))
94
95 epsilon = 1e-14
96 def loss_check3(X, w, b, y):
97     return -(tf.matmul(X, w) + b)*y + tf.math.log(1+tf.math.exp(tf.matmul(X, w) + b)+ epsilon)
98
99 def loss_check4(X, w, b, y):
100     return -(tf.matmul(X, w) + b)*y + tf.math.log(tf.clip_by_value(1+tf.math.exp(tf.matmul(X,
101     w) + b), clip_value_min=1e-15, clip_value_max=1.0))
102
103 #optimizer---sgd
104 def sgd(params, grads, lr, batch_size): #@save
105     """Minibatch stochastic gradient descent."""
106     for param, grad in zip(params, grads):
107         param.assign_sub(lr*grad/batch_size)
108     ### minibatch sgd---loss
109
110 w = tf.Variable(tf.random.normal(shape=(2, 1), mean=0, stddev=0.01),
111                 trainable=True)
112 b = tf.Variable(tf.zeros(1), trainable=True)
113
114 lr = 0.1
115 num_epochs = 10
116 net = logistic
117 loss = loss_check2
118 batch_size = 10
119 for epoch in range(num_epochs):
120     for X, y in data_iter(batch_size, X_train, y_train):
121         with tf.GradientTape() as g:
122             l = loss(X, w, b, y) # Minibatch loss in `X` and `y`
123             # Compute gradient on l with respect to `[w, b]`
124             dw, db = g.gradient(l, [w, b])
125             # Update parameters using their gradient
126             sgd([w, b], [dw, db], lr, batch_size)
127             train_l = loss(X_train, w, b, y_train)
128             print(f'epoch {epoch + 1}, loss {float(tf.reduce_mean(train_l)):f}, w {tf.reshape(w,
129             true_w.shape)}, b {b-0}')
130 # true W[2, -3.4], ture b [4.2]
131 ### sgd---loss_check2 build in loss function
132
133 w = tf.Variable(tf.random.normal(shape=(2, 1), mean=0, stddev=0.01),
134                 trainable=True)
135 b = tf.Variable(tf.zeros(1), trainable=True)
136
137 lr = 1
138 num_epochs = 1000
139 net = logistic
140 loss = loss_check2
141 batch_size = 8000
142 for epoch in range(num_epochs):
143     for X, y in data_iter(batch_size, X_train, y_train):
144         with tf.GradientTape() as g:
145             l = loss(X, w, b, y) # Minibatch loss in `X` and `y`
146             # Compute gradient on l with respect to `[w, b]`
147             dw, db = g.gradient(l, [w, b])
148             # Update parameters using their gradient
149             sgd([w, b], [dw, db], lr, batch_size)
150             train_l = loss(X_train, w, b, y_train)
151             print(f'epoch {epoch + 1}, loss {float(tf.reduce_mean(train_l)):f}, w {tf.reshape(w,

```

```

    true_w.shape)), b {b-0}')
148 # true W[2, -3.4], ture b [4.2]
149 ### minibatch sgd---loss2 build in loss function
150
151 w = tf.Variable(tf.random.normal(shape=(2, 1), mean=0, stddev=0.01),
152                 trainable=True)
153 b = tf.Variable(tf.zeros(1), trainable=True)
154
155 lr = 0.1
156 num_epochs = 5
157 net = logistic
158 loss = loss2
159 batch_size = 10
160 for epoch in range(num_epochs):
161     for X, y in data_iter(batch_size, X_train, y_train):
162         with tf.GradientTape() as g:
163             l = loss(net(X, w, b), y) # Minibatch loss in `X` and `y`
164             # Compute gradient on l with respect to `[w, b]`
165             dw, db = g.gradient(l, [w, b])
166             # Update parameters using their gradient
167             sgd([w, b], [dw, db], lr, batch_size)
168         train_l = loss(net(X_train, w, b), y_train)
169         print(f'epoch {epoch + 1}, loss {float(tf.reduce_mean(train_l)):f}, w {tf.reshape(w,
170 true_w.shape)}, b {b-0}')
171 # true W[2, -3.4], ture b [4.2]
172 ##### with penalty
173 #L2
174 def l2_penalty(w):
175     return tf.reduce_sum(tf.pow(w, 2)) / 2
176
177 ### minibatch sgd---loss
178
179 w = tf.Variable(tf.random.normal(shape=(2, 1), mean=0, stddev=0.01),
180                 trainable=True)
181 b = tf.Variable(tf.zeros(1), trainable=True)
182
183 lr = 0.1
184 num_epochs = 10
185 net = logistic
186 loss = loss_check2
187 batch_size = 10
188 lambd = 1
189 for epoch in range(num_epochs):
190     for X, y in data_iter(batch_size, X_train, y_train):
191         with tf.GradientTape() as g:
192             l = loss(X, w, b, y) + lambd * l2_penalty(w) # Minibatch loss in `X` and `y`
193             # Compute gradient on l with respect to `[w, b]`
194             dw, db = g.gradient(l, [w, b])
195             # Update parameters using their gradient
196             sgd([w, b], [dw, db], lr, batch_size)
197         train_l = loss(X_train, w, b, y_train)
198         print(f'epoch {epoch + 1}, loss {float(tf.reduce_mean(train_l)):f}, w {tf.reshape(w,
199 true_w.shape)}, b {b-0}')
200 # true W[2, -3.4], ture b [4.2]
201
202 #1.2 logistic regression manual---nba data
203 #shuffle the data
204 def data_iter(batch_size, features, labels):
205     num_examples = len(features)
206     indices = list(range(num_examples))
207     # The examples are read at random, in no particular order
208     random.shuffle(indices)
209     for i in range(0, num_examples, batch_size):
210         j = tf.constant(indices[i: min(i + batch_size, num_examples)])
211         yield tf.gather(features, j), tf.gather(labels, j)
212
213 #model
214 def net(X):
215     return tf.sigmoid(tf.matmul(X, w) + b)
216
217 #optimizer---sgd
218 def sgd(params, grads, lr, batch_size):
219     """Minibatch stochastic gradient descent."""
220     for param, grad in zip(params, grads):
221         param.assign_sub(lr*grad/batch_size)

```



```

221
222 #loss function
223 def loss(y_hat, y):
224     return tf.losses.binary_crossentropy(
225         y, y_hat, from_logits=True)
226
227 #accuracy
228 def accuracy(y_hat, y):
229     m = y.shape[0]
230     Y_prediction = np.zeros((1, m))
231     for i in range(y.shape[0]):
232         if y_hat[i] > 0.5:
233             Y_prediction[[0],[i]] = 1
234         else:
235             Y_prediction[[0],[i]] = 0
236     #ac = tf.dtypes.cast(tf.reshape(Y_prediction, [m, -1]), tf.float32)-y
237     return 1-tf.reduce_mean(tf.math.abs(tf.dtypes.cast(tf.reshape(Y_prediction, [m, -1]), tf.
238         float32)-y))
239
240 #evaluate the model on the whole test dataset
241 def evaluate_accuracy(net, data_iter):
242     metric = Accumulator(2) # No. of correct predictions, no. of predictions
243     for _, (X, y) in enumerate(data_iter):
244         metric.add(accuracy(net(X), y), 1) # use 1 to count the number of accumulation
245     return metric[0] / metric[1]
246
247 def load_array(data_arrays, batch_size, is_train=True): #@save
248     """Construct a TensorFlow data iterator."""
249     dataset = tf.data.Dataset.from_tensor_slices(data_arrays)
250     if is_train:
251         dataset = dataset.shuffle(buffer_size=1000)
252     dataset = dataset.batch(batch_size)
253     return dataset
254
255 #batch_size = 10
256 #data_iter = load_array((features, labels), batch_size)
257 #next(iter(data_iter))
258
259 class Accumulator: #@save
260     """For accumulating sums over `n` variables."""
261     def __init__(self, n):
262         self.data = [0.0] * n
263
264     def add(self, *args):
265         self.data = [a + float(b) for a, b in zip(self.data, args)]
266
267     def reset(self):
268         self.data = [0.0] * len(self.data)
269
270     def __getitem__(self, idx):
271         return self.data[idx]
272
273 #updater
274 class Updater():
275     """For updating parameters using minibatch stochastic gradient descent."""
276     def __init__(self, params, lr):
277         self.params = params
278         self.lr = lr
279
280     def __call__(self, batch_size, grads):
281         d2l.sgd(self.params, grads, self.lr, batch_size)
282
283 def train_mlp(net, train_iter, loss, updater):
284     # Sum of training loss, sum of training accuracy, no. of examples
285     metric = Accumulator(3)
286     for X, y in train_iter:
287         # Compute gradients and update parameters
288         with tf.GradientTape() as tape:
289             y_hat = net(X)
290             # Keras implementations for loss takes (labels, predictions)
291             # instead of (predictions, labels) that users might implement
292             # in this book, e.g. `cross_entropy` that we implemented above
293             l = loss(y_hat, y)
294         if isinstance(updater, tf.keras.optimizers.Optimizer):
295             params = net.trainable_variables
296             grads = tape.gradient(l, params)
297             updater.apply_gradients(zip(grads, params))
298         else:

```

```

296         updater(X.shape[0], tape.gradient(1, updater.params))
297     # Keras loss by default returns the average loss in a batch
298     l_sum = 1 * float(tf.size(y)) if isinstance(
299         loss, tf.keras.losses.Loss) else tf.reduce_sum(1)
300     metric.add(l_sum, accuracy(y_hat, y), 1)
301 # Return training loss and training accuracy
302 train_loss = metric[0] / metric[2]
303 train_accuracy = metric[1] / metric[2]
304 return train_loss, train_accuracy
305
306 #return the test accuracy & print the result step by step
307 def train_test_mlp(net, train_iter, test_iter, loss, num_epochs, updater):
308     for epoch in range(num_epochs):
309         train_metrics = train_mlp(net, train_iter, loss, updater)
310         test_acc = evaluate_accuracy(net, test_iter)
311         #animator.add(epoch + 1, train_metrics + (test_acc,))
312         train_loss, train_acc = train_metrics
313         print(f'epoch {epoch + 1}, train_loss {float(train_loss):f}, train_acc {train_acc-0},
314               test_acc {test_acc-0}')
315 #plot & test accuracy
316 def train_mlp_plot(net, train_iter, test_iter, loss, num_epochs, updater): #@save
317     """Train a model (defined in Chapter 3)."""
318     animator = Animator(xlabel='epoch', xlim=[1, num_epochs], ylim=[0, 2],
319                        legend=['train loss', 'train acc', 'test acc'])
320     for epoch in range(num_epochs):
321         train_metrics = train_with_plot(net, train_iter, loss, updater)
322         test_acc = evaluate_accuracy(net, test_iter)
323         animator.add(epoch + 1, train_metrics + (test_acc,))
324     train_loss, train_acc = train_metrics
325
326 # train loss & train accuracy
327 def train_with_plot(net, train_iter, loss, updater): #@save
328     """The training loop defined in Chapter 3."""
329     # Sum of training loss, sum of training accuracy, no. of examples
330     metric = Accumulator(3)
331     for X, y in train_iter:
332         # Compute gradients and update parameters
333         with tf.GradientTape() as tape:
334             y_hat = net(X)
335             # Keras implementations for loss takes (labels, predictions)
336             # instead of (predictions, labels) that users might implement
337             # in this book, e.g. `cross_entropy` that we implemented above
338             l = loss(y_hat, y)
339         if isinstance(updater, tf.keras.optimizers.Optimizer):
340             params = net.trainable_variables
341             grads = tape.gradient(l, params)
342             updater.apply_gradients(zip(grads, params))
343         else:
344             updater(X.shape[0], tape.gradient(1, updater.params))
345     # Keras loss by default returns the average loss in a batch
346     l_sum = 1 * float(tf.size(y)) if isinstance(
347         loss, tf.keras.losses.Loss) else tf.reduce_sum(1)
348     metric.add(l_sum, accuracy(y_hat, y), 1)
349 # Return training loss and training accuracy
350 train_loss = metric[0] / metric[2]
351 train_accuracy = metric[1] / metric[2]
352 return train_loss/35, train_accuracy
353
354 class Animator: #@save
355     """For plotting data in animation."""
356     def __init__(self, xlabel=None, ylabel=None, legend=None, xlim=None,
357                  ylim=None, xscale='linear', yscale='linear',
358                  fmts=('-', 'm--', 'g-.', 'r:'), nrows=1, ncols=1,
359                  figsize=(3.5, 2.5)):
360         # Incrementally plot multiple lines
361         if legend is None:
362             legend = []
363         d2l.use_svg_display()
364         self.fig, self.axes = d2l.plt.subplots(nrows, ncols, figsize=figsize)
365         if nrows * ncols == 1:
366             self.axes = [self.axes, ]
367         # Use a lambda function to capture arguments
368         self.config_axes = lambda: d2l.set_axes(
369             self.axes[0], xlabel, ylabel, xlim, ylim, xscale, yscale, legend)
370         self.X, self.Y, self.fmts = None, None, fmts

```

```

371     def add(self, x, y):
372         # Add multiple data points into the figure
373         if not hasattr(y, "__len__"):
374             y = [y]
375         n = len(y)
376         if not hasattr(x, "__len__"):
377             x = [x] * n
378         if not self.X:
379             self.X = [[] for _ in range(n)]
380         if not self.Y:
381             self.Y = [[] for _ in range(n)]
382         for i, (a, b) in enumerate(zip(x, y)):
383             if a is not None and b is not None:
384                 self.X[i].append(a)
385                 self.Y[i].append(b)
386         self.axes[0].cla()
387         for x, y, fmt in zip(self.X, self.Y, self.fmts):
388             self.axes[0].plot(x, y, fmt)
389         self.config_axes()
390         display.display(self.fig)
391         display.clear_output(wait=True)
392     num_classes = 1
393     num_features = 19
394     data_size = 1340
395     train_size = tf.cast(0.8*data_size, tf.int32)
396     nba = pd.read_csv("nba_logreg.csv") #shape: 1340x21
397     #fill in NA value with meanx
398     nba = nba.fillna(nba.mean())
399     # split into input and output columns
400     X, y = nba.values[:, 1:-1], nba.values[:, -1]
401     all_features = nba.iloc[:, 1:-1]
402     numeric_features = all_features.dtypes[all_features.dtypes != 'object'].index
403     all_features[numeric_features] = all_features[numeric_features].apply(
404         lambda x: (x - x.mean()) / (x.std()))
405     # After standardizing the data all means vanish, hence we can set missing
406     # values to 0
407     all_features[numeric_features] = all_features[numeric_features].fillna(0)
408     nba = all_features
409     X = nba.values
410     X = tf.convert_to_tensor(X, dtype = tf.float32)
411     y = tf.convert_to_tensor(y, dtype = tf.float32)
412     y = tf.reshape(y, [data_size,1])
413
414     # get the train/test data
415     X_train, X_test, y_train, y_test = X[0:train_size,:], X[train_size:data_size,:], y[0:
        train_size,:], y[train_size:data_size,:]
416     batch_size = 64
417     num_epochs = 10
418     lr = 0.001
419
420     #initialization
421     w = tf.Variable(tf.random.normal(shape=(19, 1), mean=0, stddev=0.01),
422         trainable=True)
423     b = tf.Variable(tf.zeros(1), trainable=True)
424
425
426     train_iter = d2l.load_array((X_train, y_train), batch_size)
427     test_iter = d2l.load_array((X_test, y_test), batch_size)
428
429     updater = d2l.Updater([w, b], lr)
430     train_test_mlp(net, train_iter, test_iter, loss, num_epochs, updater)
431     batch_size = 64
432     num_epochs = 10
433     lr = 0.001 #note: better result with 0.1
434     #initialization
435     w = tf.Variable(tf.random.normal(shape=(19, 1), mean=0, stddev=0.01),
436         trainable=True)
437     b = tf.Variable(tf.zeros(1), trainable=True)
438
439     train_iter = load_array((X_train, y_train), batch_size)
440     test_iter = load_array((X_test, y_test), batch_size)
441
442     updater = d2l.Updater([w, b], lr)
443     train_mlp_plot(net, train_iter, test_iter, loss, num_epochs, updater)
444     #note: the loss was scaled by 35 to fit in the plot
445

```

```

446
447
448
449 # regularization: L1/L2, regu parameter
450 ###adding logistic Regression---manual data
451 num_classes = 1
452 num_features = 2
453 data_size = 1000
454 train_size = tf.cast(0.8*data_size, tf.int32)
455 #data generating & processing
456 def logit_data(w, b, num_examples):
457     """Generate y = Xw + b."""
458     X = tf.zeros((num_examples, w.shape[0]))
459     X += tf.random.normal(shape=X.shape, stddev=1)
460     y = tf.matmul(X, tf.reshape(w, (-1, 1))) + b
461     #y += tf.random.normal(shape=y.shape, stddev=0.01)
462     y = tf.reshape(y, (-1, 1))
463     return X, y
464
465 true_w = tf.constant([2, -3.4])
466 true_b = 4.2
467 x, y_in = logit_data(true_w, true_b, data_size)
468 p = 1/(1+tf.exp(-y_in))
469 bernoulli_distribution = tfd.Bernoulli(probs=p)
470 X=x
471 sample = bernoulli_distribution.sample(1)
472 y = sample[0] #choose 2 dimensions out of 3 in tensor
473 y = tf.dtypes.cast(y, tf.float32)
474
475 #data processing
476 ##spilt data into training & testing training: testing = 4:1
477
478 X_train, X_test, y_train, y_test = X[0:train_size,:], X[train_size:data_size,:], y[0:
    train_size,:], y[train_size:data_size,:]
479 # Feature Matrix # Data labels
480 #X_test.shape, y_train.shape
481 batch_size = 64
482 num_epochs = 15
483 lr = 0.001
484
485 #initialization
486 w = tf.Variable(tf.random.normal(shape=(2, 1), mean=0, stddev=0.01),
487                 trainable=True)
488 b = tf.Variable(tf.zeros(1), trainable=True)
489
490
491 train_iter = d2l.load_array((X_train, y_train), batch_size)
492 test_iter = d2l.load_array((X_test, y_test), batch_size)
493
494 updater = d2l.Updater([w, b], lr)
495 train_test_mlp(net, train_iter, test_iter, loss, num_epochs, updater)
496 ###
497 batch_size = 64
498 num_epochs = 15
499 lr = 0.001
500
501 #initialization
502 w = tf.Variable(tf.random.normal(shape=(2, 1), mean=0, stddev=0.01),
503                 trainable=True)
504 b = tf.Variable(tf.zeros(1), trainable=True)
505
506 train_iter = load_array((X_train, y_train), batch_size)
507 test_iter = load_array((X_test, y_test), batch_size)
508
509 updater = d2l.Updater([w, b], lr)
510 train_mlp_plot(net, train_iter, test_iter, loss, num_epochs, updater)
511 #1.3 logistic regression Keras---nba data
512 # mlp for binary classification
513 import TensorFlow as tf
514 import pandas as pd
515 from sklearn.model_selection import train_test_split
516 from sklearn.preprocessing import LabelEncoder
517 from TensorFlow.keras import Sequential
518 from TensorFlow.keras.layers import Dense
519 nba = pd.read_csv("nba_logreg.csv") #shape: 1340x21
520 #fill in NA value with meanx

```

```

521 nba = nba.fillna(nba.mean())
522 # split into input and output columns
523 X, y = nba.values[:, 1:-1], nba.values[:, -1]
524 # ensure all data are floating point values
525 X = X.astype('float32')
526 # encode strings to integer
527 y = LabelEncoder().fit_transform(y)
528 # split into train and test datasets
529 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
530 print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)
531 # determine the number of input features
532 n_features = X_train.shape[1]
533 n_features
534 # define model
535 model = tf.keras.models.Sequential()
536 model.add(tf.keras.layers.Dense(1, activation='sigmoid', kernel_initializer='he_normal',
    input_shape=(n_features,)))
537 # compile the model
538 model.compile(optimizer='sgd', loss='binary_crossentropy', metrics=['accuracy'])
539 # fit the model
540 model.fit(X_train, y_train, epochs=100, batch_size=32, verbose=2)
541 # evaluate the model
542 loss, acc = model.evaluate(X_test, y_test, verbose=2)
543                                     #verbose:By setting verbose 0, 1 or 2 you just say how do you
    want to 'see'
544                                     #the training progress for each epoch.
545 print('Test Accuracy: %.3f' % acc)
546
547 #2.1 MLP manual
548 num_inputs, num_outputs, num_hiddens = 2, 1, 256
549 batch_size = 32
550 num_epochs = 15
551 lr = 0.01 #note: better result with 0.1
552 #initialization
553 W1 = tf.Variable(tf.random.normal(
554     shape=(num_inputs, num_hiddens), mean=0, stddev=0.01))
555 b1 = tf.Variable(tf.zeros(num_hiddens))
556 W2 = tf.Variable(tf.random.normal(
557     shape=(num_hiddens, num_outputs), mean=0, stddev=0.01))
558 b2 = tf.Variable(tf.random.normal([num_outputs], stddev=.01))
559
560 train_iter = load_array((X_train, y_train), batch_size)
561 test_iter = load_array((X_test, y_test), batch_size)
562
563 updater = d2l.Updater([W1, W2, b1, b2], lr)
564 train_mlp_plot(net, train_iter, test_iter, loss, num_epochs, updater)
565 #keras
566 # mlp for binary classification
567 import TensorFlow as tf
568 import pandas as pd
569 from sklearn.model_selection import train_test_split
570 from sklearn.preprocessing import LabelEncoder
571 from TensorFlow.keras import Sequential
572 from TensorFlow.keras.layers import Dense
573 # determine the number of input features
574 n_features = X_train.shape[1]
575 n_features
576 # define model
577 model = tf.keras.models.Sequential()
578 model.add(tf.keras.layers.Dense(256, activation='relu', kernel_initializer='he_normal',
    input_shape=(n_features,)))
579 model.add(tf.keras.layers.Dense(8, activation='relu', kernel_initializer='he_normal'))
580 model.add(tf.keras.layers.Dense(1, activation='sigmoid'))
581 # compile the model
582 model.compile(optimizer='sgd', loss='binary_crossentropy', metrics=['accuracy'])
583 # fit the model
584 model.fit(X_train, y_train, epochs=100, batch_size=32, verbose=2)
585 # evaluate the model
586 loss, acc = model.evaluate(X_test, y_test, verbose=2)
587                                     #verbose:By setting verbose 0, 1 or 2 you just say how do you
    want to 'see'
588                                     #the training progress for each epoch.
589 print('Test Accuracy: %.3f' % acc)
590 num_classes = 1
591 num_features = 19
592 data_size = 1340

```

```

593 train_size = tf.cast(0.8*data_size, tf.int32)
594 nba = pd.read_csv("nba_logreg.csv") #shape: 1340x21
595 #fill in NA value with meanx
596 nba = nba.fillna(nba.mean())
597 # split into input and output columns
598 X, y = nba.values[:, 1:-1], nba.values[:, -1]
599 all_features = nba.iloc[:, 1:-1]
600 numeric_features = all_features.dtypes[all_features.dtypes != 'object'].index
601 all_features[numeric_features] = all_features[numeric_features].apply(
602     lambda x: (x - x.mean()) / (x.std()))
603 # After standardizing the data all means vanish, hence we can set missing
604 # values to 0
605 all_features[numeric_features] = all_features[numeric_features].fillna(0)
606 nba = all_features
607 X = nba.values
608 X = tf.convert_to_tensor(X, dtype = tf.float32)
609 y = tf.convert_to_tensor(y, dtype = tf.float32)
610 y = tf.reshape(y, [data_size,1])
611
612 # get the train/test data
613 X_train, X_test, y_train, y_test = X[0:train_size,:], X[train_size:data_size,:], y[0:
    train_size:], y[train_size:data_size:]
614 #without plot
615 num_inputs, num_outputs, num_hiddens = 19, 1, 256
616 batch_size = 32
617 num_epochs = 30
618 lr = 0.01
619
620 W1 = tf.Variable(tf.random.normal(
621     shape=(num_inputs, num_hiddens), mean=0, stddev=0.01))
622 b1 = tf.Variable(tf.zeros(num_hiddens))
623 W2 = tf.Variable(tf.random.normal(
624     shape=(num_hiddens, num_outputs), mean=0, stddev=0.01))
625 b2 = tf.Variable(tf.random.normal([num_outputs], stddev=.01))
626
627 train_iter = d2l.load_array((X_train, y_train), batch_size)
628 test_iter = d2l.load_array((X_test, y_test), batch_size)
629
630 updater = d2l.Updater([W1, W2, b1, b2], lr)
631 train_test_mlp(net, train_iter, test_iter, loss, num_epochs, updater)
632 #with plot
633 num_inputs, num_outputs, num_hiddens = 19, 1, 256
634 batch_size = 32
635 num_epochs = 30
636 lr = 0.01 #note: better result with 0.1
637 #initialization
638 W1 = tf.Variable(tf.random.normal(
639     shape=(num_inputs, num_hiddens), mean=0, stddev=0.01))
640 b1 = tf.Variable(tf.zeros(num_hiddens))
641 W2 = tf.Variable(tf.random.normal(
642     shape=(num_hiddens, num_outputs), mean=0, stddev=0.01))
643 b2 = tf.Variable(tf.random.normal([num_outputs], stddev=.01))
644
645 train_iter = load_array((X_train, y_train), batch_size)
646 test_iter = load_array((X_test, y_test), batch_size)
647
648 updater = d2l.Updater([W1, W2, b1, b2], lr)
649 train_mlp_plot(net, train_iter, test_iter, loss, num_epochs, updater)
650 #2.2 MLP Keras---nba
651 nba = pd.read_csv("nba_logreg.csv") #shape: 1340x21
652 #fill in NA value with meanx
653 nba = nba.fillna(nba.mean())
654 nba
655 # split into input and output columns
656 X, y = nba.values[:, 1:-1], nba.values[:, -1]
657 # ensure all data are floating point values
658 X = X.astype('float32')
659 # encode strings to integer
660 y = LabelEncoder().fit_transform(y)
661 # split into train and test datasets
662 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33)
663 print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)
664 # determine the number of input features
665 n_features = X_train.shape[1]
666 n_features
667 # define model

```

```

668 model = tf.keras.models.Sequential()
669 model.add(tf.keras.layers.Dense(10, activation='relu', kernel_initializer='he_normal',
670     input_shape=(n_features,)))
671 model.add(tf.keras.layers.Dense(8, activation='relu', kernel_initializer='he_normal'))
672 model.add(tf.keras.layers.Dense(1, activation='sigmoid'))
673 # compile the model
674 model.compile(optimizer='sgd', loss='binary_crossentropy', metrics=['accuracy'])
675 # fit the model
676 model.fit(X_train, y_train, epochs=30, batch_size=32, verbose=2)
677 # evaluate the model
678 loss, acc = model.evaluate(X_test, y_test, verbose=2)
679                                     #verbose:By setting verbose 0, 1 or 2 you just say how do you
        want to 'see'
680                                     #the training progress for each epoch.
681 print('Test Accuracy: %.3f' % acc)
682 # make a prediction
683 row = [1,0,0.99,-0.889,0.853,0.036,0.898,-0.37,1,10,1,2,3,4,9,0.7,0.5,0.1,0.4]
684 yhat = model.predict([row])
685 print('Predicted: %.3f' % yhat)
686 #linear regression
687 %matplotlib inline
688 from d2l import TensorFlow as d2l
689 import TensorFlow as tf
690 import random
691 def synthetic_data(w, b, num_examples):
692     """Generate y = Xw + b + noise."""
693     X = tf.zeros((num_examples, w.shape[0]))
694     X += tf.random.normal(shape=X.shape)
695     y = tf.matmul(X, tf.reshape(w, (-1, 1))) + b
696     y += tf.random.normal(shape=y.shape, stddev=0.01)
697     y = tf.reshape(y, (-1, 1))
698     return X, y
699 true_w = tf.constant([2, 3.5])
700 true_b = 4
701 features, labels = synthetic_data(true_w, true_b, 1000)
702 features
703 def data_iter(batch_size, features, labels):
704     num_examples = len(features)
705     indices = list(range(num_examples))
706     # The examples are read at random, in no particular order
707     random.shuffle(indices)
708     for i in range(0, num_examples, batch_size):
709         j = tf.constant(indices[i: min(i + batch_size, num_examples)])
710         yield tf.gather(features, j), tf.gather(labels, j)
711 def data_iter(batch_size, features, labels):
712     num_examples = len(features)
713     indices = list(range(num_examples))
714     # The examples are read at random, in no particular order
715     random.shuffle(indices)
716     for i in range(0, num_examples, batch_size):
717         j = tf.constant(indices[i: min(i + batch_size, num_examples)])
718         yield tf.gather(features, j), tf.gather(labels, j)
719 w = tf.Variable(tf.random.normal(shape=(2, 1), mean=0, stddev=0.01),
720     trainable=True)
721 b = tf.Variable(tf.zeros(1), trainable=True)
722 w, b
723 def linreg(X, w, b): #@save
724     """The linear regression model."""
725     return tf.matmul(X, w) + b
726 def linreg(X, w, b): #@save
727     """The linear regression model."""
728     return tf.matmul(X, w) + b
729 def sgd(params, grads, lr, batch_size):
730     """Minibatch stochastic gradient descent."""
731     for param, grad in zip(params, grads):
732         param.assign_sub(lr*grad/batch_size)
733 w = tf.Variable(tf.random.normal(shape=(2, 1), mean=0, stddev=0.01),
734     trainable=True)
735 b = tf.Variable(tf.zeros(1), trainable=True)
736 lr = 1
737 num_epochs = 5
738 net = linreg
739 loss = squared_loss
740 batch_size = 50
741

```

```

742 for epoch in range(num_epochs):
743     for X, y in data_iter(batch_size, features, labels):
744         with tf.GradientTape() as g:
745             l = loss(net(X, w, b), y) # Minibatch loss in `X` and `y`
746             # Compute gradient on l with respect to `[w, b]`
747             dw, db = g.gradient(l, [w, b])
748             # Update parameters using their gradient
749             sgd([w, b], [dw, db], lr, batch_size)
750         train_l = loss(net(features, w, b), labels)
751         print(f'epoch {epoch + 1}, loss {float(tf.reduce_mean(train_l)):f}')
752     print(f'error in estimating w: {true_w - tf.reshape(w, true_w.shape)}')
753     print(f'error in estimating b: {true_b - b}')
754     w = tf.Variable(tf.random.normal(shape=(2, 1), mean=0, stddev=0.01),
755                     trainable=True)
756     b = tf.Variable(tf.zeros(1), trainable=True)
757
758     lr = 1
759     num_epochs = 5
760     net = linreg
761     loss = squared_loss
762     batch_size = 1000
763
764 for epoch in range(num_epochs):
765     with tf.GradientTape() as g:
766         l = loss(net(features, w, b), labels)
767         # Compute gradient on l with respect to `[w, b]`
768         dw, db = g.gradient(l, [w, b])
769         # Update parameters using their gradient
770         sgd([w, b], [dw, db], lr, batch_size)
771         train_l = loss(net(features, w, b), labels)
772         print(f'epoch {epoch + 1}, loss {float(tf.reduce_mean(train_l)):f}')
773     print(f'error in estimating w: {true_w - tf.reshape(w, true_w.shape)}')
774     print(f'error in estimating b: {true_b - b}')
775     #generate one column of ones to the left
776     x_left = tf.ones([1000, 1], tf.float32)
777     #combine the dataset with one column
778     X = tf.concat([x_left, features], 1)
779     X

```